

Timing Channels in OS Design

Comparing S3K and seL4 Time Protection

Operating Systems Comparison

February 12, 2026

Outline

- 1 S3K: Hardware-Assisted Isolation
- 2 seL4: Time Protection Abstraction
- 3 Comparison: S3K vs. seL4

S3K: Overview

- **Target:** Embedded RISC-V systems.
- **Goal:** A capability-based multicore partitioning kernel for Real-Time Systems.
- **Core Philosophy:**
 - Uses RISC-V hardware features (specifically `fence.t`) to enforce temporal isolation.
 - Minimal Trusted Computing Base (TCB) acting as a bare-metal hypervisor.
- **Key Attributes:**
 - Capability model with bounded Worst-Case Execution Time (WCET).
 - Time-driven scheduler.
 - Predictable constant-time system calls

S3K: Timing Isolation Mechanisms

S3K employs a hardware-centric approach to prevent timing interference:

The fence.t Instruction

- Used at context switches to flush core-local microarchitectural state.
- Prevents state (like branch predictors or L1 cache) from leaking data to the next process.

Scheduling & Padding

- **Major/Minor Frames:** Time is strictly partitioned.
- **Padding:** The scheduler uses worst-case padding to ensure strict deterministic process dispatch.
- **Execution:** Constant-time kernel operations prevent the kernel itself from becoming a timing channel.

S3K: Memory Hierarchy & Limitations

Scratchpad Memory (SPM) Focus:

- S3K is designed to reside entirely within the SPM (backed by L1).
- **Benefit:** Enables efficient flushing of the kernel's footprint via `fence.t`.

Scope of Protection:

- **In-Kernel:** Protected. All syscalls are non-interfering and constant-time.
- **User-Level:** Not fully protected by the kernel.
- Unlike seL4, S3K does not enforce cache coloring for user space. Processes must protect their own user-level state from side-channels in higher-level caches (L2/DRAM)

seL4: The Missing OS Abstraction

- **Problem:** Standard memory protection (spatial) is insufficient. Shared microarchitectural state (caches, TLB, branch predictors) creates timing channels.
- **Solution:** **Time Protection** as a first-class OS abstraction.
- **Definition:** Preventing interference between security domains such that execution speed in one domain is independent of another

Threat Model

Addresses both **Side Channels** (victim leaks to attacker) and **Covert Channels** (colluding sender/receiver).

seL4: Implementation Mechanisms

seL4 introduces software mechanisms to overcome hardware limitations:

① Kernel Cloning:

- Replaces static partitioning with dynamic kernel clones.
- Each domain gets a private copy of OS text/stack to prevent kernel state leakage.

② Deterministic Flushing:

- On domain switch: flush L1, TLB, BP.
- **Crucial:** Operations are padded to worst-case latency to hide the cost of the flush itself.

③ Cache Coloring:

- Partitions the Last Level Cache (LLC) by locking pages to specific cache sets.
- Prevents cross-domain interference in shared memory hierarchies.

Comparison: State Clearing Strategy

S3K (Hardware-Centric)

- Relies on RISC-V fence.t.
- **Pros:** Efficient hardware implementation; single instruction handles local state.
- **Cons:** Requires specific hardware support (RISC-V extensions); focuses mostly on core-local state.

seL4 (Software-Centric)

- Manual software flushing sequences (e.g., iterating cache lines) + Padding.
- **Pros:** Works on commodity hardware (x86/Arm) by abstracting the lack of flush instructions.
- **Cons:** Higher overhead due to manual flushing and padding requirements.

Comparison: Memory Hierarchy Protection

How they handle the cache hierarchy (L2/LLC) differs significantly:

S3K	seL4
Optimized for SPM/L1 locality.	Optimized for Deep Hierarchies (L2/LLC).
Does not provide inherent protection for shared upper-level caches (DRAM/L2).	Uses Cache Coloring to spatially partition shared caches.
<i>"User processes are responsible for protecting their own user-level state."</i>	<i>"OS enforces isolation via coloring to prevent contention."</i>

Comparison: Kernel Architecture

S3K: The Hypervisor Model

- Single kernel instance acting as a bare-metal hypervisor.
- Relies on constant-time implementation of syscalls to prevent leakage.
- **Focus:** Determinism via scheduling and instruction padding.

seL4: The Kernel Clone Model

- **Kernel Cloning:** Creates separate kernel images per domain.
- Spatially separates global kernel data (stacks, variables) to eliminate shared state channels.
- **Focus:** Spatial separation of the kernel itself to aid temporal isolation

Summary of Differences

Feature	S3K	seL4
Primary Mechanism	fence.t (RISC-V)	Software Flush + Padding
Kernel State	Single Constant-Time Kernel	Cloned Kernel Images
Shared Cache	User Responsibility	Cache Coloring (OS Managed)
Isolation Scope	Core-local (L1/SPM)	Full Hierarchy (L1-LLC)
Target	Embedded Real-Time	General Purpose / Cloud